

Evaluating LLM Reliability for Coding Adversarial Machine Learning Research Papers

Abstract—Meta-research on security and machine learning often requires systematic coding of large paper corpora, a task that has remained manual and slow. Large language models (LLMs) may help automate annotation, but their reliability for coding security papers is unknown. We evaluate LLM reliability for coding adversarial ML research by comparing four models (GPT-4o, GPT-5.2, Gemini-3, Claude Sonnet-4.5) to human expert coding on 71 papers and 9 categorical variables. All tested models reach nearly perfect intra-rater consistency (Fleiss’ $\kappa = 0.97$ – 1.00), high inter-LLM agreement ($\kappa = 0.79$ – 0.85), and substantial agreement with the human coder ($\kappa = 0.64$ – 0.71). Reliability varies by category: structured variables like domain reach almost perfect agreement ($\kappa = 0.86$), while interpretive ones like attack category are moderate ($\kappa = 0.55$). Within adversarial ML, LLMs appear usable as screening tools for structured meta-research coding, though interpretive judgment still benefits from human review.

1. Introduction

The security and privacy research community increasingly recognizes the need for critical self-reflection on research practices [1], [2], [3]. From evaluating peer review processes [4] to assessing reproducibility [5], meta-research enables evidence-based improvements to how we conduct, evaluate, and sustain scientific work. Careful meta-analysis still requires someone to code papers: methodology, evaluation practices, artifacts, threat models. That does not scale; manual coding is costly and usually keeps studies to tens of papers.

LLMs have shown decent reliability on text annotation in social science [6], [7] and medical reviews [8], with some agreement levels close to human annotators [9]. Nobody has yet tested whether they can reliably code *security* papers, where threat models, evaluation methodology, and jargon are specific. Reliable LLM coding could support large-scale meta-analysis; if not, knowing where they fail clarifies when expert coding is still needed.

Adversarial ML is a good testbed. The field has known evaluation issues [10], [11], [12], [13] (overstated robustness, weak adaptive evaluation), critiqued in manual work [14] but never in a large systematic coding study. Reliable automated coding would make that kind of meta-analysis feasible.

Contribution. We run the first evaluation of LLM reliability for coding security papers and ask: (1) How stable is each LLM across repeated codings (intra-rater)? (2) How well do LLMs match human coders? (3) Which variables are reliable or brittle under automation?

We use four LLMs (GPT-4o, GPT-5.2, Google Gemini-3, Claude Sonnet-4.5) on 71 adversarial ML papers and 9 categorical variables: research type, domain, venue, artifacts, threat model, and evaluation. Each LLM coded all papers 3 times; we compute Fleiss’ κ for intra-rater and Cohen’s κ for inter-rater agreement [15], [16].

Results are encouraging. All four models achieve nearly perfect self-consistency ($\kappa = 0.97$ – 1.00), high inter-LLM agreement ($\kappa = 0.79$ – 0.85), and substantial agreement with the human coder ($\kappa = 0.64$ – 0.71), in the range often deemed acceptable [17]. By category, structured variables (e.g., domain $\kappa = 0.86$, threat model $\kappa = 0.75$) reach high agreement, while interpretive ones (e.g., attack category $\kappa = 0.55$) are moderate.

Within adversarial ML, LLMs look useful as *screening* tools for structured variables and could cut human effort on a first pass; they still fall short of expert coding on hard methodological calls. Whether these findings generalise to other security subfields remains an open question. We report reliability numbers and release codebooks, prompts, and scripts for reuse.

2. Related Work

Recent research demonstrates that LLMs can approach human-level annotator performance on text classification tasks in social sciences [6], [7], medical reviews [8], and content moderation [9]. Gilardi et al. [6] showed that GPT-3.5 achieves close to human agreement ($\kappa = 0.7$) on tweet classification, and GPT-4 outperforms crowd workers. Ziems et al. [18] compared LLMs to human annotators across 27 datasets and found that LLMs excel on structured tasks but struggle when subjective judgment is required. Tornberg [7] showed that GPT-4 can beat experts and crowd workers on political tweets. Rathje et al. [9] achieved human-like correlation ($r = 0.59$ – 0.77) across 15 datasets and 12 languages ($n = 47,925$). The picture is mixed for systematic reviews. Khraisha et al. [19] got $\kappa = 0.65$ for GPT-4 on the screening task. However, a scoping review of 37 studies [20] described LLM “not yet ready” as a primary review tool. Gartlehner et al. [8] showed that Claude assisted extraction achieves 91% accuracy compared to 89% human based on

9,341 extractions from 63 medical studies. There have been common issues of model drift [21], prompt sensitivity [22] and extraction hallucinations. While the medical and social science fields have addressed LLM for coding papers, there remains a gap in the evaluation of LLMs for coding security papers.

Content analysis uses standard agreement metrics. Cohen’s κ [23] is used for two raters with correction chances. Fleiss’s κ [24] extends to multiple raters. Krippendorff’s α [15] is flexible across the number of coders, with commonly cited advice as $\alpha \geq 0.80$ for reliability [25]. Landis and Koch [17] suggest κ in the range of 0.61–0.80 as “substantial” and 0.81–1.0 as “almost perfect.” Neuendorf [16] suggests a $\kappa \geq 0.80$ for publication-quality work. We use Fleiss’s κ for intra-rater reliability and Cohen’s κ for inter-rater reliability. We bootstrap 95% confidence intervals. For threat-model variables T1 and T2, we compute subset κ not including N/A to avoid prevalence.

Adversarial Machine Learning has common problems in evaluation. Carlini and Wagner [10] demonstrated defensive distillation, which gave almost no robustness despite 99.5% claims. Athalye et al. [11] revealed that 7 out of 9 non-certified ICLR 2018 defences relied on obfuscated gradients. Tramer et al. [12] broke 14 more from top venues. Arp et al. [13] reviewed 30 papers over a decade and listed 10 pitfalls like data snooping, sampling bias and lab-only evaluation. All these critiques suggest what’s wrong, but not how common the pattern is. Large-scale coding of papers could reveal the pattern. Prior meta-research in security has examined peer review [4], reproducibility [5] and artefact evaluations [3]. However, most research only analyze few tens of papers due to the heavy time and resources involved in manual coding. Olszewski et al. [5] found no clear gain in artefact availability after the evaluation committee appeared in 2020 across roughly 750 ML security papers. A reliable LLM coding could save hundreds of hours in manual systematic reviews and quantify trends across entire subfields in a matter of no time.

3. Methodology

We selected papers in adversarial machine learning research based on their adoption by widely used security artefacts. This serves as a proxy for impact and visibility. We included papers if they were cited by the following artefacts: CleverHans, IBM ART, Foolbox, TextAttack, PyRIT, benchmarks (RobustBench, AutoAttack, HarmBench), or industry threat frameworks (MITRE ATLAS). The adoption of these practical artefacts ensures papers are highly relevant in security research. Papers span over a decade, from the early adversarial ML research in 2013 to 2024. They cover a range of themes from foundation attacks like FGSM [26], C&W [10], certified defences [27], robustness benchmarks [28] and recent works on LLM jailbreaking [29]. Each paper is identified by its title and arXiv ID.

We extract full text from each PDF with PyMuPDF (first 15 pages, max 50,000 characters). All raters saw the papers in the same order. For code availability variables (G5/G6),

this PDF-only input restriction may disadvantage LLMs relative to the human coder, who could access external resources (GitHub, project pages).

Each paper was coded on 9 categorical variables grouped into three themes: research characteristics (G1–G6), threat protection properties specific to attack papers (T1–T2) and practical evaluation (Q1). These coding variables are standard and are commonly seen in the literature in adversarial ML reviews like [13], [30]; our framework was inspired in particular by [30]. Table 1 includes the complete variable list and allowed values. Some categories intentionally overlap (e.g., “Attack/Defense”) to capture papers spanning multiple contributions, reflecting real-world ambiguity in the literature.

We use four LLMs and one human rater (Table 2). Based on their state-of-the-art performance in coding papers, we chose GPT-4o, GPT-5.2, Gemini-3 and Claude Sonnet-4.5 for coding our 71 papers. Each LLM was prompted separately through independent API calls with the same codebook, definition, allowed values, and a 0 temperature for consistency. We made 3 runs for each of the four models. A single human rater coded all 71 papers based on the same codebook.

Before computing agreement, we normalised all the values by trimming whitespace, hyphens, underscores, and removing None, N/A, null, and empty values, collapsing them to n/a so that formatting noise does not cause disagreement. We use two standard metrics [15], [16] to compute agreement. For models with 3 runs, we compute Fleiss’ κ [24] over the three runs with items = paper $\times$ variable. This measures the model self-consistency across different runs. For the Human coder with only 1 run, intra-rater reliability cannot be computed and is marked as N/A. For each pair of raters, we compute Cohen’s κ [23] on coding consensus. We concatenate all 9 columns into 639 labels per rater (71×9) for a global 5×5 matrix (diagonal = Fleiss’ κ for multi-run models, off-diagonal = Cohen’s κ). We also report mean κ per variable (G1–G6, T1–T2, Q1) across pairs to see which codes are stable. Interpretation follows the Landis and Koch [17] scale: $\kappa < 0.21$ (slight), 0.21–0.40 (fair), 0.41–0.60 (moderate), 0.61–0.80 (substantial), 0.81–1.00 (almost perfect).

4. Results

Figure 1 shows the full agreement matrix. Diagonal entries report intra-rater reliability (Fleiss’ κ); off-diagonal entries report inter-rater agreement (Cohen’s κ). Each κ is computed over all 639 labels (71 papers $\times$ 9 variables) for the given rater pair.

All four LLMs show nearly perfect self-consistency (intra-rater reliability): Gemini-3 achieves $\kappa = 1.000$ (identical outputs across 3 runs), Sonnet-4.5 reaches $\kappa = 0.994$, while GPT-4o and GPT-5.2 both reach $\kappa = 0.975$. Temperature-0 sampling produces highly deterministic behavior. For inter-rater reliability, all four LLMs show high inter-model agreement ($\kappa = 0.79$ – 0.85), suggesting they converge on similar interpretations. The best LLM pair

TABLE 1. CODING SCHEMA: 9 CATEGORICAL VARIABLES

Variable	Description	Allowed Values
<i>Group 1: Research Characteristics</i>		
G1	Paper type	Attack, Defense, Evaluation, Both, Attack/Defense, Attack/Evaluation, Defense/Evaluation, Attack/Defense/Evaluation
G2	Attack category (if attack)	Evasion, Poisoning, Privacy, Defense, N/A
G3	Application domain	Vision, NLP, LLM, Audio, Malware, Tabular, Cross, Cross-domain
G4	Publication venue category	ML, Security, Journal, arXiv-only
G5	Code currently available	Yes, No
G6	Code release timing	At-pub, Post-pub, Never
<i>Group 2: Threat Model (Attack Papers Only)</i>		
T1	Adversary model access	White, Black, Gray, White/Black, N/A
T2	Gradient computation required	Yes, No, N/A
<i>Group 3: Practical Evaluation</i>		
Q1	Evaluation on real-world system	Yes, No, Partial, N/A

TABLE 2. RATERS AND RUN CONFIGURATION

Rater	Model/Source	Runs	Consensus
GPT-4o	OpenAI API	3	Majority vote
GPT-5.2	OpenAI API	3	Majority vote
Gemini-3	Google Gemini 3.0 Pro	3	Majority vote
Sonnet-4.5	claude-sonnet-4-5-20250929	3	Majority vote
Human	Domain expert	1	N/A




Figure 1. Global agreement matrix. Diagonal (dark green): intra-rater Fleiss’ κ for all LLMs (3 runs each); N/A for the single-run Human coder. Off-diagonal: inter-rater Cohen’s κ . Color scale: red (poor) $\rightarrow$ green (excellent).

is GPT-5.2 $\leftrightarrow$ Gemini-3 ($\kappa = 0.846$), while Sonnet-4.5 achieves excellent agreement with other LLMs ($\kappa = 0.79$ – 0.81). All LLMs achieve substantial agreement with the human coder: Gemini-3 ($\kappa = 0.707$), GPT-5.2 ($\kappa = 0.675$),

GPT-4o ($\kappa = 0.654$), and Sonnet-4.5 ($\kappa = 0.643$), all in the “substantial” range that Landis and Koch [17] consider acceptable for published work. Table 3 shows mean κ across all 10 inter-rater pairs, ranked by reliability; per-variable heatmaps appear in Appendix A.

TABLE 3. PER-VARIABLE AGREEMENT (MEAN κ ACROSS ALL PAIRS WITH 95% CI)

Var	Description	Mean κ	95% CI	Interp.
G3	Application domain	.856	[.802, .901]	Almost perfect
T1	Threat model access	.748	[.674, .819]	Substantial
G5	Code available	.742	[.658, .821]	Substantial
T2	Gradient required	.736	[.644, .822]	Substantial
G6	Code release timing	.699	[.614, .780]	Substantial
Q1	Real-world evaluation	.614	[.508, .716]	Moderate
G4	Publication venue	.611	[.482, .734]	Moderate
G1	Paper type	.589	[.469, .706]	Moderate
G2	Attack category	.548	[.418, .674]	Moderate

Structured variables perform best: application domain (G3) achieves almost perfect agreement ($\kappa = 0.86$), while threat-model access (T1), gradient requirement (T2), code availability (G5), and release timing (G6) all reach substantial agreement ($\kappa = 0.70$ – 0.75). These have well-defined categories with clear textual signals (e.g., “white-box,” “black-box,” “vision,” “NLP”). Interpretive variables are harder: publication venue (G4), real-world evaluation (Q1), paper type (G1), and attack category (G2) achieve moderate agreement ($\kappa = 0.55$ – 0.61). These require judgment. For example, is a paper that proposes an attack and then defends against it “Attack,” “Defense,” or “Both”? Is a poisoning attack also “privacy” if it extracts training data? Disagreements arise from ambiguous boundaries. The best pairs per variable reach substantial to almost perfect agreement: G3 (domain) achieves $\kappa = 0.87$ (GPT-5.2 $\leftrightarrow$ Sonnet-4.5), T2 (gradient) reaches $\kappa = 0.90$ (GPT-5.2 $\leftrightarrow$ Gemini-3), and G5 (code available) hits $\kappa = 0.89$ (GPT-5.2 $\leftrightarrow$ Gemini-3). Even G2 (attack category) achieves $\kappa = 0.79$ between GPT-5.2 and Sonnet-4.5.

One caveat concerns threat-model variables: while

T1/T2 appear to reach substantial agreement in Table 3, subset κ excluding N/A reveals much lower LLM-human agreement ($\kappa = 0.18\text{--}0.33$). LLMs marked 56–70% of papers as N/A for T1/T2, while the human marked only 41–45%, indicating these variables are less reliable than global metrics suggest and benefit from N/A-driven prevalence inflation (see Limitations).

5. Discussion

Some variables reach a high agreement, and some don’t. This gap reveals what LLMs can and can’t extract from a text. Application domain (G3; $k = 0.86$) has a high agreement because it maps to concrete, repeated lexical signals. Papers testing on Imagenet, CIFAR-10, or MNIST mention these datasets in their abstracts and methodology. Vision papers include terms like “image”, “CNN”, and “pixel”. NLP papers mention “text”, “embedding”, or specific models, for example, BERT. These bag-of-words features correlate strongly with the ground truth. LLMs trained on millions of texts have learned these associations. The task here seems more pattern matching than free interpretation.

Paper type (G1, $k = 0.59$) and attack category (G2, $k = 0.55$) obtain a moderate alignment due to ambiguity. A paper on Autoattack is an attack method and an evaluation framework (it benchmarks defences). A human coder can interpret this as “attack” because they know it’s primarily an attack contribution by the author’s framing, and making a wise judgment. LLMs see the paper superficially but struggle with the metalevel question of the author’s intent. The codebook allows overlapping categories like Attack/Evaluation, but this only shifts the ambiguity, not resolves it.

Real-world evaluation (Q1, $k = 0.61$) lies in the middle. LLMs can detect explicit mentions of real-world deployment practices such as Google Cloud Vision API, Tesla Autopilot, and GPT-4 through pattern matching. However, they struggle with the boundary between “partial deployment” and “no deployment at all”. A paper that tests on a realistic traffic sign dataset is not deployed on real-world systems, but is closer to real-world conditions than CIFAR-10. A human expert can apply their background knowledge and infer how close the adversarial ML paper is to deployment. On the other hand, LLMs lack a contextual model and rely on surface-level cues, misclassifying surface-level cues as “Yes” if the paper mentions real-world relevance.

Threat model variables T1 and T2 perform at a substantial agreement ($k = 0.75, 0.74$), but the results might be misleading. When we compute k excluding N/A, agreement drops to $k = 0.18 - 0.33$ for human-LLM pairs. The high global k is inflated based on papers which are not attacks (hence, LLMs are supposed to respond N/A). LLMs overassign N/A (56-70% of papers) compared to the human expert (41-45%). This represents a conservative strategy in coding, where we mark a paper as N/A when unsure whether it is an attack or not. Feinstein and Cicchetti [31] call this a “prevalence induced k ”, where most items fall into one category (N/A). LLMs can identify papers that clearly describe gradient-based white-box attacks, but struggle with

boundary cases like defences that include ablation attacks, or evaluation papers that propose new attacks as a baseline.

The lesson is that LLMs can reliably automate extraction that reduces to lexical pattern matching over subjective categories. Application domain, code availability (“Github”, “code”) and gradient use (∇ symbol, “backpropagation”) are extractable because they have stable surface forms. Coding tasks that have conceptual ambiguity, where applying background knowledge or inferring the author’s intent remains unreliable. This aligns with findings in social science annotation [18] where sentiment and stance detection (lexical) outperform intent classification (subjective).

For meta-research in adversarial ML—and potentially other areas of security—this suggests a practical division of work. Use LLMs to precode structured metadata (venue, domain, artefacts, threat model access) at scale, significantly reducing human effort. Reserve human coding for interpretive variables (paper type, primary contribution, evaluation rigour) where disagreements reflect genuine ambiguity in the literature. Our correlation results with humans ($k = 0.64\text{--}0.71$) suggest that LLM output can serve as an important first pass screening, not the ground truth. The right workflow is LLM-assisted coding with human adjudication in automation. Moreover, the screening automation might also help researchers decide if their direction or choice of paper is feasible for the research question they are examining, before they start spending their valuable time coding the paper. This might help accelerate the pace of meta research on discovering new patterns based on existing papers, as “rough coding” and testing proof of concept become easier.

These findings carry broader implications for adversarial ML meta-research. Our 71-paper corpus serves as a proof of concept demonstrating that reliable LLM coding could enable systematic studies of evaluation practices, artefact availability, and threat-model reporting across hundreds of papers, the kind of quantitative meta-analysis that manual coding has historically precluded. Researchers could, for instance, track whether artefact-evaluation policies [32] correlate with actual code availability over time, or examine whether vulnerabilities such as “obfuscated gradients” [11] persist in defenses. Our reliability numbers provide initial evidence that such large-scale studies are feasible for structured variables, at least within adversarial ML; further validation across other security subfields is needed.

6. Limitations

One of the major limitations of our work arises from resource constraints. First, we only use one expert coder due to the large number of papers and coding criteria. A single human coder may have bias, which can be mitigated by averaging or majority voting among three human raters. Prior work suggests human-human κ for similar tasks is 0.6–0.8 [15]; our LLM-human $\kappa = 0.71$ is strong, but this should be validated with multiple humans. Second, we were able to test only a few state-of-the-art models due to budget limitations for API calls. We would need to test in future with other models which show promise: Opus

4.6, o3 reasoning. As these models keep evolving, future models might show different results than models in early 2026. Third, we only used a single one-shot prompt and did not optimise. Per-model tuning or a few-shot examples could improve reliability. Fourth, the intra-rater consistency of models across three runs is not particularly significant and is largely a consequence of keeping the temperature to 0. Finally, the models might be influenced by their training data on the papers we included. The choice of top adversarial ML papers may not generalise to other security fields (systems security, cryptography).

7. Conclusion

We evaluated LLM reliability for coding adversarial ML papers on 9 categorical variables and found that all models achieve nearly perfect self-consistency ($\kappa = 0.97$ – 1.00) and substantial agreement with a human expert ($\kappa = 0.64$ – 0.71), with reliability varying by variable type: structured codes like domain reach almost perfect agreement ($\kappa = 0.86$), while interpretive codes like attack category are moderate ($\kappa = 0.55$ – 0.61). These results indicate that LLMs are usable as screening tools for structured meta-research coding in adversarial ML, though they do not yet replace expert coding for nuanced methodological judgment and generalisability to other security subfields remains to be tested. All supplementary materials including codebook, prompts, analysis scripts, raw CSV outputs, and per-model coding results will be made publicly available upon acceptance.

8. Ethics Considerations

We evaluate LLM coding of public academic papers. No human subjects; all papers are published and openly or institutionally accessible. No privacy, consent, or harm issues. The goal is to improve how the security community does meta-research.

9. LLM Usage Considerations

LLMs were both *evaluated* (as coders) and *used* (literature search, running experiments, analysis). We treat their coding output as data and validate it against human coding. LLMs helped with literature search, paper writing on LaTeX, running experiments, and debugging. We used Claude code for writing experiment codes and verified everything manually. We release prompts, scripts, and raw outputs for replication.

References

- [1] “MetaCRISP 2026: 1st Workshop on Metascience and Critical Reflections in Security & Privacy,” <https://metacrisp.org/>, 2026, co-located with IEEE S&P 2026.
- [2] C. Rossow, C. J. Dietrich, C. Grier, C. Kreibich, V. Paxson, N. Pohlmann, H. Bos, and M. van Steen, “Prudent practices for designing malware experiments: Status quo and outlook,” in *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, 2012, pp. 65–79.
- [3] A. Baset and T. Denning, “A data-driven reflection on 36 years of security and privacy research,” in *Proceedings of the 12th USENIX Workshop on Cyber Security Experimentation and Test (CSET)*, 2019.
- [4] V. Le Pochat, “Reflecting on research practices: Encouraging the security research community to engage in meta-research,” *Communications of the ACM*, 2024, opinion, Computing Profession. Posted Mar. 27 2024.
- [5] D. Olszewski, A. Lu, C. Stillman, K. Warren, C. Kitroser, A. Pascual, D. Ukirde, K. R. B. Butler, and P. Traynor, ““get in researchers; we’re measuring reproducibility”: A reproducibility study of machine learning papers in tier 1 security conferences,” in *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2023, pp. 2349–2363.
- [6] F. Gilardi, M. Alizadeh, and M. Kubli, “ChatGPT outperforms crowd workers for text-annotation tasks,” *Proceedings of the National Academy of Sciences*, vol. 120, no. 30, p. e2305016120, 2023.
- [7] P. Törnberg, “ChatGPT-4 outperforms experts and crowd workers in annotating political twitter messages with zero-shot learning,” *arXiv preprint arXiv:2304.06588*, 2024.
- [8] G. Gartlehner, S. Kugley, K. Crotty *et al.*, “Artificial intelligence–assisted data extraction with a large language model: A study within reviews,” *Annals of Internal Medicine*, vol. 178, pp. 1763–1771, 2025, epub 4 November 2025. doi: 10.7326/ANNALS-25-00739.
- [9] S. Rathje, D.-M. Mirea, I. Sucholutsky, R. Marjeh, C. Robertson, and J. J. Van Bavel, “GPT is an effective tool for multilingual psychological text analysis,” *Proceedings of the National Academy of Sciences*, vol. 121, no. 8, p. e2308950121, 2024.
- [10] N. Carlini and D. Wagner, “Towards evaluating the robustness of neural networks,” in *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, 2017, pp. 39–57.
- [11] A. Athalye, N. Carlini, and D. Wagner, “Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples,” in *Proceedings of the International Conference on Machine Learning (ICML)*, 2018, pp. 274–283.
- [12] F. Tramèr, N. Carlini, W. Brendel, and A. Madry, “On adaptive attacks to adversarial example defenses,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 1633–1645.
- [13] D. Arp, E. Quiring, C. Wressnegger, and K. Rieck, “Dos and don’ts of machine learning in computer security,” in *Proceedings of the USENIX Security Symposium*, 2022, pp. 3971–3988, distinguished Paper Award.
- [14] N. Carlini, A. Athalye, N. Papernot, W. Brendel, J. Rauber, D. Tsipras, I. Goodfellow, A. Madry, and A. Kurakin, “On evaluating adversarial robustness,” *arXiv preprint arXiv:1902.06705*, 2019.
- [15] K. Krippendorff, *Content Analysis: An Introduction to Its Methodology*, 2nd ed. Sage Publications, 2004.
- [16] K. A. Neuendorf, *The Content Analysis Guidebook*, 2nd ed. Sage Publications, 2017.
- [17] J. R. Landis and G. G. Koch, “The measurement of observer agreement for categorical data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.
- [18] C. Ziems, W. Held, O. Shaikh, J. Chen, Z. Zhang, and D. Yang, “Can large language models transform computational social science?” *Computational Linguistics*, vol. 50, no. 1, pp. 237–291, 2024.
- [19] Q. Khraisha, S. Put, J. Kappenberg, A. Warraitch, and K. Hadfield, “Can large language models replace humans in systematic reviews? evaluating GPT-4’s efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages,” *Research Synthesis Methods*, 2024, early view. doi: 10.1002/jrsm.1715.
- [20] J.-L. Lieberum, M. Toews, M.-I. Metzendorf, F. Heilmeyer, W. Siemens, C. Haverkamp, D. Böhringer, J. J. Meerpohl, and A. Eisele-Metzger, “Large language models for conducting systematic reviews: On the rise, but not yet ready for use—a scoping review,” *Journal of Clinical Epidemiology*, vol. 181, p. 111746, 2025.

- [21] L. Chen, M. Zaharia, and J. Zou, “How is ChatGPT’s behavior changing over time?” *arXiv preprint arXiv:2307.09009*, 2023.
- [22] M. Sclar, Y. Choi, Y. Tsvetkov, and A. Suhr, “Quantifying language models’ sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting,” *arXiv preprint arXiv:2310.11324*, 2023.
- [23] J. Cohen, “A coefficient of agreement for nominal scales,” *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960.
- [24] J. L. Fleiss, “Measuring nominal scale agreement among many raters,” *Psychological Bulletin*, vol. 76, no. 5, pp. 378–382, 1971.
- [25] A. F. Hayes and K. Krippendorff, “Answering the call for a standard reliability measure for coding data,” *Communication Methods and Measures*, vol. 1, no. 1, pp. 77–89, 2007.
- [26] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2015, arXiv:1412.6572.
- [27] J. M. Cohen, E. Rosenfeld, and J. Z. Kolter, “Certified adversarial robustness via randomized smoothing,” in *Proceedings of the International Conference on Machine Learning (ICML)*, 2019, pp. 1310–1320, arXiv:1902.02918.
- [28] F. Croce and M. Hein, “Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks,” in *Proceedings of the International Conference on Machine Learning (ICML)*, 2020, arXiv:2003.01690.
- [29] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [30] G. Apruzzese, H. S. Anderson, S. Dambra, D. Freeman, F. Pierazzi, and K. A. Roundy, ““real attackers don’t compute gradients”: Bridging the gap between adversarial ML research and practice,” in *Proceedings of the IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, 2023, pp. 339–364.
- [31] A. R. Feinstein and D. V. Cicchetti, “High agreement but low kappa: I. the problems of two paradoxes,” *Journal of Clinical Epidemiology*, vol. 43, no. 6, pp. 543–549, 1990.
- [32] “USENIX Security ’25 Call for Artifacts,” <https://www.usenix.org/conference/usenixsecurity25/call-for-artifacts>, 2025, accessed: 2026-02-12.

Appendix

Appendix: Supplementary Material and Coding Prompt

Figure 2 shows Cohen’s κ for each variable across all rater pairs. Diagonal entries show intra-rater reliability (Fleiss’ κ) for all four LLMs (each with 3 runs); the Human coder diagonal is N/A (single run). Variables G3 (domain), G5 (code available), T1 (threat model), and T2 (gradient) show consistent high agreement across most LLM pairs. Variables G1 (paper type), G2 (attack category), and Q1 (real-world evaluation) show moderate agreement, reflecting their interpretive nature.

All LLMs received identical system prompts with temperature set to 0.0 for deterministic outputs. Each paper was provided as extracted PDF text (first 15 pages, max 50,000 characters) with title and arXiv ID. The 71 papers were selected based on artifact-adoption criteria detailed in Section 3.1.

All supplementary materials will be made publicly available upon acceptance, including raw CSV outputs from each model run (3 runs per LLM), human expert codings, the complete paper list with arXiv IDs and PDF links, analysis scripts for computing κ statistics, and API interaction code for each model.

Coding Prompt

You are coding academic papers for adversarial ML research. For each paper, assign values to exactly 9 variables. Use ONLY the allowed values listed below.

```
## CODING VARIABLES (STRICT VALUES)

### G1: Paper Type
ALLOWED: "Attack", "Defense", "Evaluation", "Both", "Attack/Defense", "Attack/Evaluation", "Defense/Evaluation", "Attack/Defense/Evaluation"

### G2: Threat Category
ALLOWED: "Evasion", "Poisoning", "Privacy", "Defense", "N/A"

### G3: Domain
ALLOWED: "Vision", "NLP", "LLM", "Audio", "Malware", "Tabular", "Cross", "Cross-domain"

### G4: Publication Venue
ALLOWED: "ML", "Security", "Journal", "arXiv-only"

### G5: Code Available
ALLOWED: "Yes", "No"

### G6: Code Release Timing
ALLOWED: "At-pub", "Post-pub", "Never"

### T1: Model Access Level (attacks only)
ALLOWED: "White", "Black", "Gray", "White/Black", "N/A"

### T2: Gradient Required (attacks only)
ALLOWED: "Yes", "No", "N/A"

### Q1: Real-World Evaluation
ALLOWED: "Yes", "No", "Partial", "N/A"

## OUTPUT FORMAT
Return ONLY valid JSON (no markdown):
{
  "G1": "<value>", "G2": "<value>",
  "G3": "<value>", "G4": "<value>",
  "G5": "<value>", "G6": "<value>",
  "T1": "<value>", "T2": "<value>",
  "Q1": "<value>",
  "reasoning": "Brief explanation"
}

CRITICAL: Use ONLY the allowed values. Do NOT create new values.
```

Per-Column Agreement (Cohen's Kappa)



Figure 2. Per-variable agreement heatmaps showing pairwise κ for each coding variable (G1–G6: research characteristics; T1–T2: threat model; Q1: evaluation). Diagonal entries show intra-rater Fleiss' κ for all LLMs; Human is N/A (single run). Color scale ranges from red (poor) to green (excellent).